

# Quant GANs: Deep Generation of Financial Time Series

## Machine Learning for Finance

Laure Michaud - Tim Houdayé


September 18, 2022




- 1 Architecture of neural networks
  - Dilated Causal Convolutional layer
  - Advantage of the architecture
  - A rigorous definition of the TCNs for the first time
  - Advantages / Disadvantages of TCNs
- 2 Generative Adversarial Networks
  - The GAN objective
- 3 Numerical Application
  - The models
  - The results
  - Our implementation
- 4 Conclusion

# Dilated Causal Convolutional layer

- *Causal Convolution* are convolutions where the output only depends on the past and the present values not the future values.



- In order to construct *Temporal Convolutional Networks*, we will use the dilated ones where  $K$  denoted the kernel size and  $D$  the dilatation factors.



(a)  $K = D = 1$



(b)  $K = 2, D = 1$



(c)  $K = D = 2$

# Advantage of the architecture

- Possibility to look back **without increasing** the number of weights. Indeed, as we want to deal with long-range dependencies, increasing the filter could be a solution to capture more information but it will require more weights. This is why a **dilatation factor** is introduced.
- One of the drawbacks of this structure is that we cannot see values between  $x_{t-D}$  and  $x_t$ . To overcome this issue, the authors stack convolutions with different dilatation factors which will lead to the *Vanilla TCNs*. As for the *MLP*, *Vanilla TCNs* are obtained by composition of dilated causal convolution layers with activation functions.
- The article gives for the **first time a rigorous mathematical definition** of the *TCNs* which constitutes an important step in deep learning.


# A rigorous definition of the TCNs for the first time

## Temporal Convolutional Networks

A temporal convolutional network with  $L$  hidden layers is a function  $f : \mathbb{R}^{N_0 \times T_0} \rightarrow \mathbb{R}^{N_{L+1} \times T_L}$  such that:

$$f(X, \theta) = w \circ \psi_L \circ \dots \circ \psi_1(X)$$

where :  $\psi_l : \mathbb{R}^{N_{l-1} \times T_{l-1}} \rightarrow \mathbb{R}^{N_l \times T_l}$  is a block module (lipschitz function),  
 $w : \mathbb{R}^{N_L \times T_L} \rightarrow \mathbb{R}^{N_{L+1} \times T_L}$  a  $1 \times 1$  convolutional layer and  $T_0, L, N_0, \dots, N_{L+1}, S_l \in \mathbb{N}$  such that  $T_l := T_{l-1} - S_l$  leading to  $T_L - T_0 = \sum_{l=1}^L S_l$ .



# Advantages/Disadvantages of TCNs

## Advantages

- **Parallelism**
- **Flexible RFS:** a *TCN* can change its RFS in multiple ways. For example, stacking more dilated causal convolutional layers, using larger dilatation factors, or increasing the filter size. Easy to adapt to different domains.

## Disadvantages

- **Data storage during the evaluation**
- **Potential parameter change for a transfer of domain**

# The Generative Adversarial Networks

While using GANs,  $(\mathbb{R}^{N_z}, \mathcal{B}(\mathbb{R}^{N_z}))$  and  $(\mathbb{R}^{N_x}, \mathcal{B}(\mathbb{R}^{N_x}))$  are called the **latent** and the **data measure space**, respectively. The variable  $Z$  represented the **noise prior** and  $X$  the **target random variable**.

## Generator

Let  $g : \mathbb{R}^{N_z} \times \Theta^{(g)} \rightarrow \mathbb{R}^{N_x}$  be a network with a parameter space  $\Theta^{(g)}$ . The random variable  $X$ , defined as follows is called the **generated random variable**. The network  $g$  is called the **generator**, and  $X_\theta$  the **generated random variable with parameter  $\theta$** .

$$\begin{aligned} X : \Omega \times \Theta^{(g)} &\rightarrow \mathbb{R}^{N_x} \\ (\omega, \theta) &\mapsto g_\theta(Z(\omega)) \end{aligned}$$

## Discriminator

Let  $d : \mathbb{R}^{N_x} \times \Theta^{(d)} \rightarrow \mathbb{R}$  be a network with parameters  $\eta \in \Theta^{(d)}$  and  $\sigma : \mathbb{R} \mapsto [0,1] : x \mapsto \frac{1}{1+e^{-x}}$  be sigmoid function. A function  $d : \mathbb{R}^{N_x} \times \Theta^{(d)} \mapsto [0,1]$  defined by  $d : (x, \eta) \mapsto \sigma \circ d_\eta(x)$  is called a **discriminator**.

# The GAN objective

The goal of the generator  $g$  is to generate a sample ( $X$ ) such that the discriminator is unable to distinguish whether the sample were generated or sampled from the target. Therefore, we obtain the following *GAN objective* :

$$\min_{\theta \in \Theta(g)} \max_{\eta \in \Theta(d)} \mathcal{L}(\theta, \eta)$$

where :

$$\mathcal{L}(\theta, \eta) := \mathbb{E}[\log(d_\eta(X))] + \mathbb{E}[\log(1 - d_\eta(g_\theta(Z)))] = \mathbb{E}[\log(d_\eta(X))] + \mathbb{E}[\log(1 - d_\eta(\tilde{X}_\theta))].$$

- **Pros:** efficient replication of a random variable distribution - No need of an explicit density function - flexibility.
- **Cons:** hard to know when stopping the training of the generator and discriminator - no explicit representation of the generator's density - hard to invert generated model to get back to the latent variables.



Replication of the log return of the S&P 500 from 05/01/2009 to 31/12/2018 using :

- **Models:** Pure TCN, C-SVNN and GARCH(1,1)
- **Evaluation metrics:** Earth mover distance, DY metrics, ACF score, Leveraged effect score
- **Objective:** Capture all the stylized effects of logarithmic returns and try to generate "fake" returns whose characteristics are closest to the real ones.

# The models

## Pure TCN

They modeled the log returns directly using a TCN with a RFS  $T^{(g)} = 127$  as the generator function. The return process is given by :

$$R_{t,\theta} = g_{\theta}(Z_{t-(T^{(g)}-1):t})$$

for the three-dimensional noise prior  $Z_t \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(\mathbf{0}, \mathbb{I})$ , ( $N_Z=3$ ), and a RFS  $T^{(g)} = 127$  as the generator function.

## Constrained log return NP C-SVNN with drift

$$R_{t,\theta} = \sigma_{t,\theta}\epsilon_{t,\theta} + \mu_{t,\theta}$$

with volatility NP  $\sigma_{t,\theta}$ , drift NP  $\mu_{t,\theta}$ , and innovation NP  $\epsilon_{t,\theta}$  constrained to being i.i.d.  $\mathcal{N}(0,1)$ -distributed. The latent process is still  $Z_t \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(\mathbf{0}, \mathbb{I})$  for  $N_Z = 3$ . The innovation process takes the form  $\epsilon_{t,\theta} = Z_{t,1}$  for any  $t \in \mathbb{Z}$ .

## GARCH(1,1) with constant drift

This model is well known to model the returns of financial time series. The model is defined as follows :

$$\begin{aligned}R_{t,\theta} &= \xi_t + \mu \\ \xi_t &= \sigma_t \epsilon_t \\ \sigma_t^2 &= \omega + \alpha \xi_{t-1}^2 + \beta \omega_{t-1}^2 \\ \epsilon_t &\stackrel{\text{i.i.d}}{\sim} \mathcal{N}(0, 1)\end{aligned}$$

for  $\mu \in \mathbb{R}$ ,  $\omega > 0$ , and  $\alpha, \beta \in [0,1]$  such that  $\alpha + \beta < 1$ , and the parameters vector  $\theta = (\omega, \alpha, \beta, \mu)$ .

The authors gave the following results regarding their numerical application.


- The C-SVNN model gives results that are above the 2 others regarding the evaluation metrics we give.
- The pure TCN model is slightly comparable to the C-SVNN as their results are close
- The GARCH(1,1) model is the worst one. The model achieves to well capture the autocorrelation but fails to capture the exactness and the leverage effect.

# Our implementation

## Packages used

- Classical packages : numpy, pandas, random, matplotlib
- Collect of data from Yahoo Finance : yfinance
- Statistical treatments : scipy, statsmodel
- NN packages : torch

We follow the preprocessing steps mentioned in the article and the architecture of NN described in the Annex B.



# Our results : Lambert transforms

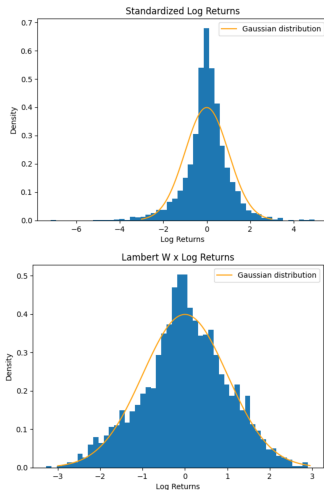


Figure: Log returns and gaussian density function

# Our results : generated log path





Figure: Generated logpaths (left : 5 paths - right : 50 paths)

# Our results : densities of S&P500




Figure: Comparison of generated and historical densities of the S&P500



- *GAN* variation which utilizes temporal convolutional networks (*TCNs*) aiming at capturing long-range dependencies like volatility clusters.
- One of the first solid work applying *GANs* to financial time series generations.
- The objective was to approximate a realistic asset price simulator by using a neural network, data-driven concept.
- Competitive results, which can be used to approximate financial time series.
- Lack of a unified metric to quantify the goodness of fit of two datasets, but overall the findings were a solid step in developing data-driven models in finance.

# References



BAI, S., ZICO, J., KORN, R., AND KOLTUN, V.

An Empirical Evaluation of Generic Convolutional and Recurrent Networks for Sequence Modeling.

*arXiv:1803.01271* (2018).



GEORG, M.

Lambert W random variables - a new family of generalized skewed distributions with applications to risk estimation.

*The Annals of Applied Statistics* (2010).



GOODFELLOW, I., POUGET-ABADIE, J., MIRZA, M., XU, B., WARDE-FARLEY, D., OZAIR, S., COURVILLE, A., AND BENGIO, Y.

Generative adversarial nets.

*<http://papers.nips.cc/paper/5423-generative-adversarial-nets.pdf>* (2014).



WIESE, M., KNOBLOCH, R., KORN, R., AND KRETSCHLMER, P.

Quant GANs: Deep Generation of Financial Time Series.